

Section 01 — LLM API Architecture

Token Economics — Groq

Speed

- Tokens per second (T/Sec)

Price

- Price per 1M token
 - input price → price on tokens user sends to the model
 - output price → price on tokens generated by the model

Rate Limits

Max usage of the model

- Tokens per Minute (TPM) → Max tokens the model can process per minute
- Request per Minute (RPM) → Max number of requests per minute
- Audio Seconds per Hour → Max audio seconds can be processed in an hour.

Context Window & Completion Length

Context Window → Max total amount of text that the model can process in one request (user prompt + model response).

Max Completion Tokens → The maximum output length the model can generate.
